

Prudhvi Raj Medikonduri

San Jose, CA • (669) 499-6998 • prudhvi.medikonduri@gmail.com • linkedin.com/in/prudhvi-raj-medikonduri

Software Engineer with 5 years of experience designing and building production backend systems in Java and Spring Boot at Goldman Sachs and Credit Karma/Intuit. Built event-driven pipelines on AWS SQS, REST microservices, and data pipelines on GCP. On the AI side, shipped an Amazon Textract document processing pipeline, built an NLP-to-SQL system, and designed the data layer for a GenAI agent built on MCP. Have owned work across the full stack from backend services and infrastructure through React frontends.

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, TypeScript, SQL, PowerShell, BeanShell

Backend: Spring Boot, Spring Security, Spring Data JPA, REST APIs, Microservices, Event-Driven Architecture, Node.js

Frontend: React, JavaScript, HTML/CSS

Cloud & Infra: AWS (EC2, S3, Lambda, SQS, RDS, IAM) | GCP (Cloud Run Jobs, Cloud Asset API, GCS, Workload Identity) | Terraform, Docker, Kubernetes, GitHub Actions, CircleCI, Jenkins, HashiCorp Vault

Databases: PostgreSQL, MySQL, Oracle SQL, MongoDB, BigQuery; schema design, query optimization, stored procedures

Messaging: Apache Kafka, RabbitMQ, AWS SQS

AI / ML: Amazon Textract, NLP, FastAPI, MCP (Model Context Protocol), data pipeline design

Testing: JUnit 5, Mockito, TestContainers, integration testing

Auth & Access: SailPoint IIQ, Okta, OAuth 2.0, OIDC, SAML, RBAC, SOX compliance

PROFESSIONAL EXPERIENCE

Senior Engineer | **Credit Karma / Intuit** | San Jose, CA

Nov 2024 – Present

- Built a **Java/Spring Boot event-driven service** combining REST APIs with an AWS SQS pipeline to handle access provisioning during Credit Karma's Google Workspace migration to Intuit. The service published access request events to an enterprise SQS queue, processed acknowledgments from downstream systems, and synced entitlement records to keep provisioning state consistent.
- Designed and built an **org-wide IAM data pipeline** in Java, Docker, and CircleCI running as scheduled Cloud Run Jobs on GCP. Resolved IAM bindings via the Cloud Asset API, joined them with entitlement exports, and published structured access catalog reports to GCS. Replaced a fully manual quarterly audit process.
- Architected the **data layer for an MCP-based GenAI agent** deployed internally at Credit Karma. On-call engineers used it to answer GCP access questions in natural language instead of parsing IAM binding files manually. Owned data modeling, catalog schema design, and the ingestion pipeline that fed the agent over MCP.
- Automated **GCP IAM provisioning** through a Terraform Enterprise GitOps workflow enforcing least-privilege controls with a full git audit trail. Integrated HashiCorp Vault into non-prod pipelines to remove hardcoded credentials across the org.
- Built a **Slack notification service** in Java that pushed real-time access request alerts to requesters, approvers, and downstream teams, with a self-service Forward Approvals flow for managers. Cut access request turnaround from 2-3 days down to a few hours, a **90%+ SLA improvement**.
- Designed and delivered **JML lifecycle automation** in Java covering provisioning on hire, role-change entitlement updates, and automated revocation on exit. Replaced a manual ops process and brought **3,000 identities across 10 applications** under automated governance in the GCP org.
- Built a custom **SailPoint IIQ connector** from scratch in Java/BeanShell with full account aggregation, provisioning, and role mapping against a SQL database. Also shipped an Okta-to-SailPoint identity correlation rule, a PagerDuty on-call access integration, and a Google Workspace enforcement script covering **5,000 groups**.

Engineering Lead | **Diadem Capital** | New York, NY

May 2024 – Nov 2024

- Designed and built a **startup-investor matching engine** in Java/Spring Boot backed by PostgreSQL, evaluating 100+ variables per startup to generate ranked recommendations across 5 investor types (VC, family office, corporate VC, venture debt, angel). Each investor type had a dedicated React dashboard with role-specific deal flow views.
- Led a 9-person engineering team building a FinTech matching platform for **1,700 investors**. Owned the microservices architecture, React/JavaScript dashboards, REST API design, and AWS-based CI/CD pipelines that cut deployment effort by **40%**.
- Built Salesforce CRM and document processor integrations via **REST webhooks and batch jobs**, automating the investor onboarding workflow and eliminating manual back-office data operations.

Senior Software Engineer | **Goldman Sachs** | Bengaluru, India

Jul 2019 – Sep 2022

- Built an **automated financial reconciliation system** in Java/Spring Boot processing millions of financial records daily from multiple upstream trading platforms. Replaced a largely manual workflow and cut reconciliation effort by **75%**. The configurable rule-based engine handled discrepancy detection, exception management, and full audit trail generation.
- Designed Oracle SQL schemas for reconciliation state management and reduced **query execution time by 40%** through targeted indexing, stored procedure rewrites, and query plan analysis.
- Maintained **85%+ test coverage** across 200+ JUnit 5, Mockito, and TestContainers tests. Handled L3 production support for high-priority financial incidents under strict SLAs.

PROJECTS

Automatic Grader | Python, FastAPI, Amazon Textract, React, MongoDB, AWS Lambda

- AI-powered exam grading platform used in production by SJSU professors. Professors uploaded batches of handwritten exams through a React frontend. An Amazon Textract pipeline extracted content, graded submissions against a rubric, and generated per-student feedback reports automatically. Backend in Python/FastAPI with MongoDB, deployed on AWS Lambda and S3.

SQL-GPT | Python, NLP, React, custom SQL parser

- Natural-language-to-SQL system where users queried a database in plain English through a React interface. The backend combined NLP models with a custom schema-aware SQL parser. Achieved **20% faster query processing** and **30% lower response time** compared to the baseline rule-based approach.

EDUCATION

M.S. Software Engineering | San Jose State University, San Jose, CA

Aug 2022 – May 2024

B.Tech. Computer Science and Engineering | V.R. Siddhartha Engineering College, India

Jul 2015 – May 2019